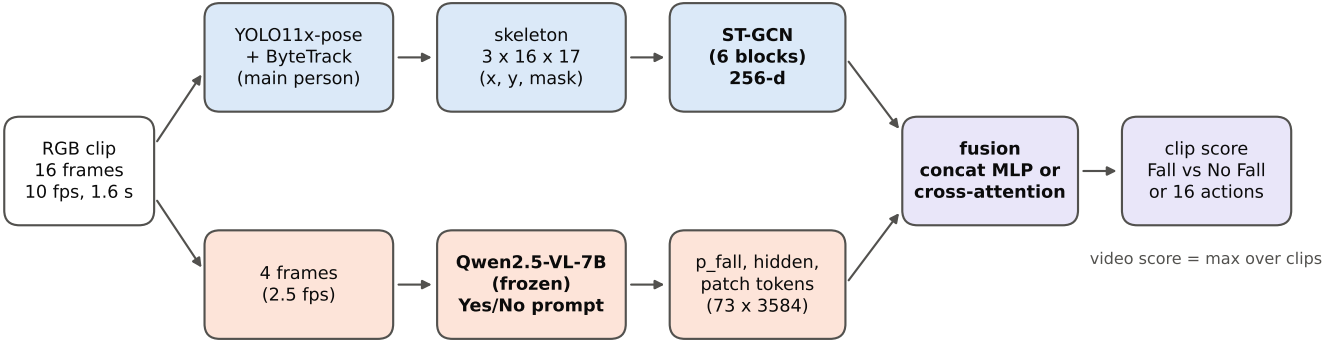
RGB clip
16 frames
10 fps, 1.6 s
YOLO11x-pose
+ ByteTrack
(main person)
skeleton
3 x 16 x 17
(x, y, mask)
ST-GCN
(6 blocks)
256-d
4 frames
(2.5 fps)
Qwen2.5-VL-7B
(frozen)
Yes/No prompt
p_fall, hidden,
patch tokens
(73 x 3584)
fusion
concat MLP or
cross-attention
clip score
Fall vs No Fall
or 16 actions
video score = max over clips